



Operating Systems

Synchronization Tools-Part3

Seyyed Ahmad Javadi

sajavadi@aut.ac.ir

Fall 2023

Copyright Notice

Slides are based on the slides of the main **textbook**.

Silberschatz

<https://www.os-book.com/OS10/slide-dir/index.html>



Mutex Locks

- ***Previous solutions are complicated*** and generally inaccessible to application programmers.
- OS designers build software tools to solve critical section problem
- Simplest is **mutex** lock
 - Boolean variable indicating if lock is available or not
 - In fact, the term mutex is short for mutual exclusion.



Mutex Locks

■ Protect a critical section by

- First **acquire()** a lock
- Then **release()** the lock

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;  
}
```

```
release() {  
    available = true;  
}
```

■ Calls to acquire() and release() must be *atomic*

- Usually implemented via hardware atomic instructions such as compare-and-swap.



Solution to CS Problem Using Mutex Locks

```
while (true) {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
}
```

Requirement	Yes/No
Mutual Exclusion	
Progress	
Bounded waiting	

- But this solution requires **busy waiting**.
 - This lock therefore called a **spinlock**.
 - Is this a disadvantage all the time?

Busy Waiting: Advantage or Disadvantage?

- Advantage of Spinlocks: no context switch is required
 - When a process must wait on a lock
 - A context switch may take considerable time

- When we prefer **spinlocks (on multi core systems)**?
 - If a lock is to be held **for a short duration**
 - One thread can “spin” on one processing core while another thread performs its critical section on another core.



Busy Waiting: Advantage or Disadvantage?

On modern multicore computing systems, spinlocks are
widely used in many operating systems.

Semaphore

- Synchronization tool that provides more sophisticated ways (than Mutex locks) for processes to synchronize their activities.
- Semaphore S – *integer variable*.
- Can only be accessed via two indivisible (atomic) operations
 - wait() and signal()
 - ▶ Originally called P() and V()



Semaphore (Cont.)

- Definition of the wait() operation

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

- Definition of the signal() operation

```
signal(S) {  
    S++;  
}
```

Semaphore (Cont.)

- **Counting semaphore** – integer value can range over an unrestricted domain.
- **Binary semaphore** – integer value can range only between 0 and 1.
 - Same as a **mutex lock**
- Can implement a counting semaphore S as a binary semaphore.
- **With semaphores we can solve various synchronization problems.**



Semaphore Usage Example

- Solution to the CS Problem
 - Create a semaphore “**mutex**” initialized to 1

```
wait(mutex) ;
```

CS

```
signal(mutex) ;
```

Requirement	Yes/No
Mutual Exclusion	
Progress	
Bounded waiting	



Semaphore Usage Example (Cont.)

- Consider P_1 and P_2 that with two statements S_1 and S_2 and the requirement that **S_1 to happen before S_2**

- Create a semaphore “**synch**” initialized to 0

P1 :

S_1 ;

signal (synch) ;

P2 :

wait (synch) ;

S_2 ;



Semaphore Implementation

- Must guarantee that no two processes can execute the **wait()** and **signal()** on the same semaphore at the same time.
- Thus, the implementation becomes the critical section problem where the **wait** and **signal** code are placed in the critical section.



Semaphore Implementation

- Could now have **busy waiting** in critical section implementation
 - But implementation code is short
 - Little busy waiting if critical section rarely occupied
- Note that applications may spend lots of time in critical sections and therefore **this is not a good solution.**



Semaphore Implementation with no Busy waiting

- With each semaphore there is an associated waiting queue.

- Each entry in a waiting queue has two data items:
 - Value (of type integer)
 - Pointer to next record in the list.



Semaphore Implementation with no Busy waiting

- Two operations:
 - **Block**
 - ▶ Place the process invoking the operation on the appropriate waiting queue
 - **Wakeup**
 - ▶ Remove one of processes in the waiting queue and place it in the ready queue



Implementation with no Busy waiting (cont.)

- Waiting queue

```
typedef struct {  
    int value;  
    struct process *list;  
} semaphore;
```

Implementation with no Busy waiting (Cont.)

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}  
  
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```

Problems with Semaphores

- Incorrect use of semaphore operations:
 - `signal(mutex) ... wait(mutex)`
 - `wait(mutex) ... wait(mutex)`
 - Omitting of `wait(mutex)` and/or `signal(mutex)`
- These – and others – are examples of what can occur when semaphores and other synchronization tools are used incorrectly.



Summary of What Are Not Covered

- **Monitors**
- **Condition Variables**
- **We also skip chapter 7 slides**
 - <https://www.os-book.com/OS10/slide-dir/index.html>

